

## FINE GRAINED CLASSIFICATION

### Data Loading and Pre-Processing

Instructions for Loading data:

1. This image file is unzipped to a folder ("/content/Deep\_Learning\_coursework/bird") within the content drive - This makes the process faster than unzipping in the drive
2. The images will now be split to Test(50%), Train(40%) and Validation(10%) using stratified split. The images are moved to the folders according to this directory:

/content

```
└─ Deep_Learning_coursework
.....└─ bird
.....└─ CUB_200_2011
.....└─ images
.....└─ test_images
.....└─ val_images
.....└─ train_images
```

4. The images are loaded using the `image_dataset_from_directory` which is an inbuilt pre-processing function from Keras.
5. While extracting the data, the size of the image (height = 256, width = 256) and batch size (32) are adjusted

### Implemented approach for fine-grained bird species classification:

- Transfer Learning tried with EfficientNetB4, ResNet, fixed on EfficientNetV2-S
- Fine Tuning done by unfreezing the last 50 layers so the model can learn from the dataset
- Model Accuracy currently at 90.62%, with precision at 93.52 and recall at 90.28
- Confusion Matrix shows that most images are perfectly classified

### Other Approches which can improve model performance:

- DenseNet121, ViT, GANs may produce identical or improved results since they are perform well with datasets of a small size
- Ensemble or model stacking
- Model has to be implemented with the larger dataset
- Other optimizers like AdaGrad, SGD (Stochastic Gradient Descent), AdamW
- Hyperparameter tuning for optimizing learning rate, batch size, dropout rate

### Code Implementation

#### Data Preprocessing:

- Images resized to 256×256, batch size 32
- Applied data augmentation (flipping, rotation, zoom, contrast adjustment)
- Used AUTOTUNE to prefetch batches for efficient training
- One-hot encoding applied to class labels (20 classes in total)

**Model Selection & Training:**

- Used EfficientNet, pre-trained on ImageNet
- Added GlobalAveragePooling2D, Dropout (0.3), and Dense (softmax) output layer for 20 classes
- Compiled with Adam optimizer and categorical cross-entropy loss
- Fine-tuned by unfreezing the last 50 layers

**Training the Model:**

- Trained on train\_ds, validated on val\_ds
- Metrics tracked: accuracy


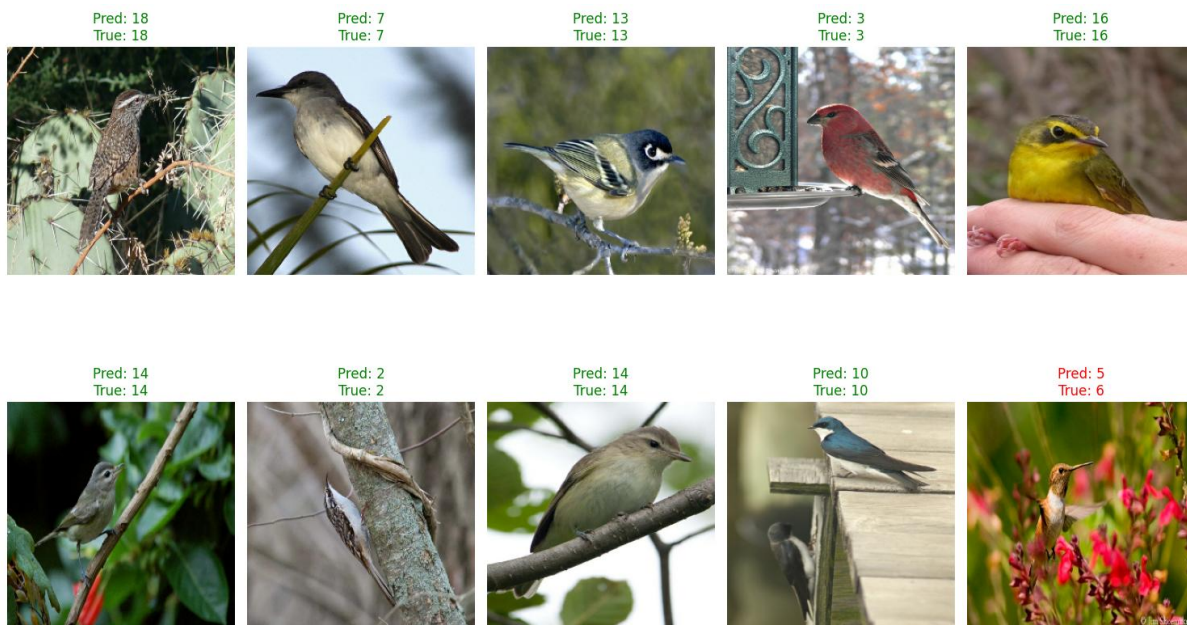
**Testing Data:**

- Test images loaded and preprocessed similarly to training data
- Model predictions displayed with actual images
- Accuracy, Precision and Recall used for performance analysis
- Confusion matrix plotted for performance analysis

**Performance Improvement:**

1. Hyperparameter tuning (highest feasibility, major impact).
2. Ensemble/Stacking (if computational resources allow).
3. Trying AdamW/SGD with tuning (helps convergence and generalization).
4. DenseNet121, ViT, GANs (ViT is great but may require larger datasets or fine-tuning).
5. Larger Dataset (ideal but dataset availability is a constraint).

**Model Results:**



▼ Performance Metrics

```
[ ] # Computing and printing performance metrics
accuracy = accuracy_score(true_classes, predicted_classes)
print(f"Accuracy: {accuracy:.4f}")

precision = precision_score(true_classes, predicted_classes, average='macro', zero_division = 1)
print(f"Precision: {precision:.4f}")

recall = recall_score(true_classes, predicted_classes, average='macro')
print(f"Recall: {recall:.4f}")
```

↗ Accuracy: 0.9862  
Precision: 0.9352  
Recall: 0.9028